

Lab #1 - Gapminder Dataset

Econ 224

August 28th, 2018

Installing Required Packages

Welcome to the first lab of Econ 224! Today we'll be giving you a crash course in two R packages that we'll be using throughout the semester: `dplyr` and `ggplot2`. Before we can get started, you'll need to install both of these packages. A quick way to install both of them at once, along with several other packages that may come in handy later, is `install.packages('tidyverse')`. Note that you only need to do this *once*. The dataset we'll work with today is also available as an R package called `gapminder`. Make sure that you have both `tidyverse` and `gapminder` installed before continuing.

The Gapminder Dataset

Our next step is to load both `tidyverse`, which contains `dplyr` and `ggplot2`, and `gapminder`, which contains the data we'll be analyzing today:

```
library(tidyverse)
library(gapminder)
```

Exercise #1

Now that you've loaded `gapminder`, use the command `?gapminder` to view the R help file for this dataset and read the documentation you find there and answer the following questions:

- What information does this dataset contain?
- How many rows and columns does it have?
- What are the names of each of the columns, and what information does each contain?
- What is the source of the dataset?

Solution to Exercise # 1

Answer Goes Here

What is a tibble?

Let's see what happens if we display the `gapminder` dataset:

```
gapminder
```

```
## # A tibble: 1,704 x 6
##   country    continent year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

If you're used to working with dataframes in R, this may surprise you. Rather than filling up the screen with lots of useless information, R shows us a helpful summary of the information contained in `gapminder`. This is because `gapminder` is *not* a dataframe; it's a *tibble*. For the moment, all you need to know about tibbles is that they are souped up versions of R dataframes that are designed to work seamlessly with `dplyr`. (If you want to learn more, see Chapter 7 of *R for Data Science*) But what is `dplyr` in the first place?

What is dplyr?

The `dplyr` package provides a number of powerful but easy-to-use tools for data manipulation in R. We'll be making heavy use of this package throughout the semester. Rather than trying to explain everything in advance, let's just dive straight in.

Filter Rows with filter

Let's run the following command in R and see what happens:

```
gapminder %>% filter(year == 2007)
```

```
## # A tibble: 142 x 6
##   country    continent year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      2007   43.8  31889923    975.
## 2 Albania     Europe      2007   76.4   3600523   5937.
## 3 Algeria     Africa      2007   72.3  33333216   6223.
## 4 Angola      Africa      2007   42.7  12420476   4797.
## 5 Argentina   Americas    2007   75.3  40301927  12779.
## 6 Australia   Oceania     2007   81.2  20434176  34435.
## 7 Austria     Europe      2007   79.8   8199783   36126.
## 8 Bahrain     Asia        2007   75.6    708573   29796.
## 9 Bangladesh  Asia        2007   64.1 150448339   1391.
## 10 Belgium    Europe      2007   79.4  10392226  33693.
## # ... with 132 more rows
```

Compare the results of running this command to what we got when we typed `gapminder` into the console above. Rather than displaying the whole dataset, now R is only showing us the 142 rows for which the column `year` has a value of 2007.

So how does this work? The `%>%` symbol is called a *pipe*. Pipes play very nicely with `dplyr` and make our code very easy to understand. The tibble `gapminder` is being piped into the function `filter()`. The argument `year == 2007` tells `filter()` that it should find all the rows such that the logical condition `year == 2007` is TRUE.

Oh no! Have we accidentally deleted all of the other rows of `gapminder`? Nope: we haven't made any changes to `gapminder` at all. If you don't believe me try entering `gapminder` at the console. All that this command does is *display* a subset of `gapminder`. If we wanted to store the result of running this command, we'd need to assign it to a variable, for example

```
gapminder2007 <- gapminder %>% filter(year == 2007)
gapminder2007
```

```
## # A tibble: 142 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int> <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      2007  43.8  31889923    975.
## 2 Albania    Europe    2007  76.4   3600523   5937.
## 3 Algeria    Africa    2007  72.3  33333216   6223.
## 4 Angola     Africa    2007  42.7  12420476   4797.
## 5 Argentina  Americas  2007  75.3  40301927  12779.
## 6 Australia  Oceania   2007  81.2  20434176  34435.
## 7 Austria    Europe    2007  79.8   8199783   36126.
## 8 Bahrain    Asia      2007  75.6    708573   29796.
## 9 Bangladesh Asia      2007  64.1 150448339   1391.
## 10 Belgium   Europe    2007  79.4  10392226   33693.
## # ... with 132 more rows
```

Exercise #2

1. Explain the difference between `x = 3` and `x == 3` in R.
2. Use `filter` to choose the subset of `gapminder` for which `year` is 2002.
3. If you instead try to choose the subset with `year` equal to 2005, something will go wrong. Try it and explain what happens and why.
4. Store the data for Asian countries in a tibble called `gapminder_asia`. Display this tibble.

Solution to Exercise #2

1. The first assigns the value 3 to the variable `x`; the second tests whether `x` is equal to 3 and returns either TRUE or FALSE.
2. Use the following code:

```
gapminder %>% filter(year == 2002)
```

```
## # A tibble: 142 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int> <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      2002  42.1  25268405    727.
## 2 Albania    Europe    2002  75.7   3508512   4604.
## 3 Algeria    Africa    2002  71.0  31287142   5288.
```

```
## 4 Angola      Africa      2002      41.0 10866106      2773.
## 5 Argentina   Americas    2002      74.3 38331121      8798.
## 6 Australia   Oceania     2002      80.4 19546792     30688.
## 7 Austria     Europe      2002      79.0  8148312     32418.
## 8 Bahrain     Asia        2002      74.8   656397     23404.
## 9 Bangladesh  Asia        2002      62.0 135656790      1136.
## 10 Belgium    Europe      2002      78.3 10311970     30486.
## # ... with 132 more rows
```

3. If you go back to the help file for `gapminder` you'll see that it only contains data for every fifth year. The year 2005 isn't in our dataset so `dplyr` will display an empty tibble:

```
gapminder %>% filter(year == 2005)
```

```
## # A tibble: 0 x 6
## #   ... with 6 variables: country <fct>, continent <fct>, year <int>,
## #     lifeExp <dbl>, pop <int>, gdpPercap <dbl>
```

4. Use the following code:

```
gapminder_asia <- gapminder %>% filter(continent == 'Asia')
gapminder_asia
```

```
## # A tibble: 396 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 386 more rows
```

Filtering two variables

We can use `filter` to subset on two or more variables. For example, here we display data for the US in 2007:

```
gapminder %>% filter(year == 2007, country == 'United States')
```

```
## # A tibble: 1 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 United States Americas    2007   78.2 301139947   42952.
```

Exercise # 3

1. When I displayed data for the US in 2007, I put quotes around `United States` but not around `year`. Explain why.
2. Which country had the higher life expectancy in 1977: Ireland or Brazil? Which had the higher GDP per capita?

Solution to Exercise # 3

1. This is because `year` contains numeric data while `country` contains character data, aka string data.
2. From the results of the following code, we see that Ireland had both a higher life expectancy and GDP per capita.

```
gapminder %>% filter(year == 1977, country == 'Ireland')
```

```
## # A tibble: 1 x 6
##   country continent  year lifeExp      pop gdpPercap
##   <fct>   <fct>     <int>  <dbl>   <int>    <dbl>
## 1 Ireland Europe    1977   72.0 3271900   11151.
```

```
gapminder %>% filter(year == 1977, country == 'Brazil')
```

```
## # A tibble: 1 x 6
##   country continent  year lifeExp      pop gdpPercap
##   <fct>   <fct>     <int>  <dbl>   <int>    <dbl>
## 1 Brazil  Americas    1977   61.5 114313951    6660.
```