

Lab #5 - Predictive Regression I

Econ 224

September 11th, 2018

Introduction

This lab provides a crash course on least squares regression in R. In the interest of time we'll work with a very simple, but somewhat boring, dataset that requires very little explanation. In our next lab and on the problem set you'll use what you've learned here to look at much more interesting examples!

The `mtcars` Dataset

The built-in R dataset `mtcars` contains information on 32 models of automobile from 1973-74 as reported in *Motor Trend Magazine*. For more information on the variables, see the R help file `?mtcars`. Note that `mtcars` is a dataframe rather than a tibble. Just to keep things simple I won't convert it to a tibble. But don't worry: everything I demonstrate in this tutorial will work just as well with a tibble as with a dataframe. A tibble *is* a dataframe even though a dataframe is not a tibble. (C.f. a square is a rectangle, but a rectangle is not a square.) Here are the first few rows of the `mtcars`:

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Our goal will be to predict `mpg` (fuel efficiency in miles/gallon) using the other variables such as `cyl` (# of cylinders), `disp` (engine displacement in cubic inches), `hp` (horsepower), and `wt` (weight in thousands of pounds).

The `lm` Command

The command for least squares regression in R is `lm` which stands for *linear model*. The basic syntax is as follows: `lm([Y variable] ~ [1st predictor] + ... + [pth predictor], [dataframe])`. For example, to predict `mpg` using `disp` and `hp` we would run the command

```
lm(mpg ~ disp, mtcars)
```

Call:

```
lm(formula = mpg ~ disp, data = mtcars)
```

Coefficients:

(Intercept)	disp
29.59985	-0.04122

Exercise #1

Carry out a regression predicting mpg using disp, hp, cyl and wt

```
lm(mpg ~ disp + hp + cyl + wt, mtcars)
```

Call:

```
lm(formula = mpg ~ disp + hp + cyl + wt, data = mtcars)
```

Coefficients:

(Intercept)	disp	hp	cyl	wt
40.82854	0.01160	-0.02054	-1.29332	-3.85390

Getting More Information from lm

If we simply run `lm` as above, R will display only the estimated regression coefficients: $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p$ along with the command used to run the regression: `Call`. To get more information, we need to *store* the results of our regression.

```
reg1 <- lm(mpg ~ disp + hp, mtcars)
```

If you run the preceding line of code in the R console, it won't produce any output. But if you check your R environment after running it, you'll see a new List object: `reg1`. To see what's inside this list, we can use the command `str`:

```
str(reg1)
```

List of 12

```
$ coefficients : Named num [1:3] 30.7359 -0.0303 -0.0248
..- attr(*, "names")= chr [1:3] "(Intercept)" "disp" "hp"
$ residuals    : Named num [1:32] -2.15 -2.15 -2.35 1.23 3.24 ...
..- attr(*, "names")= chr [1:32] "Mazda RX4" "Mazda RX4 Wag" "Datsun 710" "Hornet 4 Drive" ...
$ effects      : Named num [1:32] -113.65 -28.44 5.8 1.1 3.01 ...
..- attr(*, "names")= chr [1:32] "(Intercept)" "disp" "hp" "" ...
$ rank         : int 3
$ fitted.values: Named num [1:32] 23.1 23.1 25.1 20.2 15.5 ...
..- attr(*, "names")= chr [1:32] "Mazda RX4" "Mazda RX4 Wag" "Datsun 710" "Hornet 4 Drive" ...
$ assign       : int [1:3] 0 1 2
$ qr           :List of 5
..$ qr        : num [1:32, 1:3] -5.657 0.177 0.177 0.177 0.177 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. ..$ : chr [1:32] "Mazda RX4" "Mazda RX4 Wag" "Datsun 710" "Hornet 4 Drive" ...
.. .. ..$ : chr [1:3] "(Intercept)" "disp" "hp"
.. ..- attr(*, "assign")= int [1:3] 0 1 2
..$ qraux: num [1:3] 1.18 1.09 1.01
..$ pivot: int [1:3] 1 2 3
..$ tol   : num 1e-07
..$ rank  : int 3
..- attr(*, "class")= chr "qr"
$ df.residual : int 29
```

```

$ xlevels      : Named list()
$ call        : language lm(formula = mpg ~ disp + hp, data = mtcars)
$ terms       :Classes 'terms', 'formula' language mpg ~ disp + hp
.. ..- attr(*, "variables")= language list(mpg, disp, hp)
.. ..- attr(*, "factors")= int [1:3, 1:2] 0 1 0 0 0 1
.. ..- attr(*, "dimnames")=List of 2
.. ..- attr(*, "term.labels")= chr [1:2] "disp" "hp"
.. ..- attr(*, "order")= int [1:2] 1 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. ..- attr(*, "predvars")= language list(mpg, disp, hp)
.. ..- attr(*, "dataClasses")= Named chr [1:3] "numeric" "numeric" "numeric"
.. ..- attr(*, "names")= chr [1:3] "mpg" "disp" "hp"
$ model       :'data.frame': 32 obs. of 3 variables:
..$ mpg : num [1:32] 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
..$ disp: num [1:32] 160 160 108 258 360 ...
..$ hp : num [1:32] 110 110 93 110 175 105 245 62 95 123 ...
..- attr(*, "terms")=Classes 'terms', 'formula' language mpg ~ disp + hp
.. ..- attr(*, "variables")= language list(mpg, disp, hp)
.. ..- attr(*, "factors")= int [1:3, 1:2] 0 1 0 0 0 1
.. ..- attr(*, "dimnames")=List of 2
.. ..- attr(*, "term.labels")= chr [1:2] "disp" "hp"
.. ..- attr(*, "order")= int [1:2] 1 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. ..- attr(*, "predvars")= language list(mpg, disp, hp)
.. ..- attr(*, "dataClasses")= Named chr [1:3] "numeric" "numeric" "numeric"
.. ..- attr(*, "names")= chr [1:3] "mpg" "disp" "hp"
- attr(*, "class")= chr "lm"

```

Don't panic: you don't need to know what all of these list elements are. The important thing to understand is that `lm` returns a *list* from which we can extract important information about the regression we have run. To extract the regression coefficient estimates, we use `coef`

```
coef(reg1)
```

```

(Intercept)      disp      hp
30.73590425 -0.03034628 -0.02484008

```

To extract the regression residuals, we use `resid`

```
resid(reg1)
```

```

      Mazda RX4      Mazda RX4 Wag      Datsun 710
-2.1480911      -2.1480911      -2.3483788
Hornet 4 Drive  Hornet Sportabout      Valiant

```

1.2258440	3.2357695	-3.1997835
Duster 360	Merc 240D	Merc 230
0.5745752	-0.3440204	-1.3033408
Merc 280	Merc 280C	Merc 450SE
-3.3945383	-4.7945383	-1.4951865
Merc 450SL	Merc 450SLC	Cadillac Fleetwood
-0.5951865	-2.6951865	-0.9202450
Lincoln Continental	Chrysler Imperial	Fiat 128
-1.0359995	3.0296761	5.6917931
Honda Civic	Toyota Corolla	Toyota Corona
3.2529931	6.9363213	-3.1818286
Dodge Challenger	AMC Javelin	Camaro Z28
-1.8597761	-2.5846240	-0.7288875
Pontiac Firebird	Fiat X1-9	Porsche 914-2
4.9496206	0.6008970	1.1752002
Lotus Europa	Ford Pantera L	Ferrari Dino
5.3569558	2.2734203	-2.2886799
Maserati Bora	Volvo 142E	
1.7197522	-2.9564359	

and to extract the *fitted values* i.e. the predicted values of Y , we use `fitted.values`

```
fitted.values(reg1)
```

Mazda RX4	Mazda RX4 Wag	Datsun 710
23.14809	23.14809	25.14838
Hornet 4 Drive	Hornet Sportabout	Valiant
20.17416	15.46423	21.29978
Duster 360	Merc 240D	Merc 230
13.72542	24.74402	24.10334
Merc 280	Merc 280C	Merc 450SE
22.59454	22.59454	17.89519
Merc 450SL	Merc 450SLC	Cadillac Fleetwood
17.89519	17.89519	11.32025
Lincoln Continental	Chrysler Imperial	Fiat 128
11.43600	11.67032	26.70821
Honda Civic	Toyota Corolla	Toyota Corona
27.14701	26.96368	24.68183
Dodge Challenger	AMC Javelin	Camaro Z28
17.35978	17.78462	14.02889
Pontiac Firebird	Fiat X1-9	Porsche 914-2
14.25038	26.69910	24.82480
Lotus Europa	Ford Pantera L	Ferrari Dino
25.04304	13.52658	21.98868
Maserati Bora	Volvo 142E	
13.28025	24.35644	


Exercise # 2

1. Plot a histogram of the residuals from `reg1` using `ggplot` with a bin width of 1.25. Is there anything noteworthy about this plot?
2. Calculate the residuals “by hand” by subtracting the fitted values from `reg1` from the column `mpg` in `mtcars`. Use the R function `all.equal` to check that this gives the same result as `resid`.

Solution to Exercise #2

1. There seems to be some right skewness in the residuals.

```
library(ggplot2)
ggplot() +
  geom_histogram(aes(x = resid(reg1)), binwidth = 1.25)
```



2. They give exactly the same result:

```
all.equal(resid(reg1), mtcars$mpg - fitted.values(reg1))
```

```
[1] TRUE
```

Summarizing Regression Output

To view the usual summary of regression output, we use the `summary` command:

```
summary(reg1)
```

Call:

```
lm(formula = mpg ~ disp + hp, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.7945	-2.3036	-0.8246	1.8582	6.9363

```

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 30.735904   1.331566  23.083 < 2e-16 ***
disp        -0.030346   0.007405  -4.098 0.000306 ***
hp          -0.024840   0.013385  -1.856 0.073679 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.127 on 29 degrees of freedom
Multiple R-squared:  0.7482,    Adjusted R-squared:  0.7309
F-statistic: 43.09 on 2 and 29 DF,  p-value: 2.062e-09

```

Among other things, `summary` shows us the coefficient estimates and associated standard errors. It also displays the t-value (Estimate / SE) and associated p-value for a test of the null hypothesis $H_0: \beta = 0$ versus $H_1: \beta \neq 0$. Farther down in the output, `summary` provides the residual standard error and R-squared. It turns out the `summary` command *itself* returns a list. In particular,

```
str(summary(reg1))
```

```

List of 11
 $ call      : language lm(formula = mpg ~ disp + hp, data = mtcars)
 $ terms     :Classes 'terms', 'formula' language mpg ~ disp + hp
 .. ..- attr(*, "variables")= language list(mpg, disp, hp)
 .. ..- attr(*, "factors")= int [1:3, 1:2] 0 1 0 0 0 1
 .. ..- attr(*, "dimnames")=List of 2
 .. ..$ : chr [1:3] "mpg" "disp" "hp"
 .. ..$ : chr [1:2] "disp" "hp"
 .. ..- attr(*, "term.labels")= chr [1:2] "disp" "hp"
 .. ..- attr(*, "order")= int [1:2] 1 1
 .. ..- attr(*, "intercept")= int 1
 .. ..- attr(*, "response")= int 1
 .. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
 .. ..- attr(*, "predvars")= language list(mpg, disp, hp)
 .. ..- attr(*, "dataClasses")= Named chr [1:3] "numeric" "numeric" "numeric"
 .. ..- attr(*, "names")= chr [1:3] "mpg" "disp" "hp"
 $ residuals : Named num [1:32] -2.15 -2.15 -2.35 1.23 3.24 ...
 ..- attr(*, "names")= chr [1:32] "Mazda RX4" "Mazda RX4 Wag" "Datsun 710" "Hornet 4 Drive" ...
 $ coefficients : num [1:3, 1:4] 30.7359 -0.0303 -0.0248 1.3316 0.0074 ...
 ..- attr(*, "dimnames")=List of 2
 .. ..$ : chr [1:3] "(Intercept)" "disp" "hp"
 .. ..$ : chr [1:4] "Estimate" "Std. Error" "t value" "Pr(>|t|)"
 $ aliased     : Named logi [1:3] FALSE FALSE FALSE
 ..- attr(*, "names")= chr [1:3] "(Intercept)" "disp" "hp"
 $ sigma       : num 3.13
 $ df          : int [1:3] 3 29 3
 $ r.squared   : num 0.748
 $ adj.r.squared: num 0.731
 $ fstatistic  : Named num [1:3] 43.1 2 29
 ..- attr(*, "names")= chr [1:3] "value" "numdf" "dendf"
 $ cov.unscaled : num [1:3, 1:3] 1.81e-01 -1.18e-04 -8.38e-04 -1.18e-04 5.61e-06 ...
 ..- attr(*, "dimnames")=List of 2
 .. ..$ : chr [1:3] "(Intercept)" "disp" "hp"
 .. ..$ : chr [1:3] "(Intercept)" "disp" "hp"
 - attr(*, "class")= chr "summary.lm"

```

This fact can come in handy when you want to *extract* some of the values from the regression summary table to use for some other purpose. For example, we can display *only* the R-squared as follows: We could do this as follows:

```
summary(reg1)$r.squared
```

```
[1] 0.7482402
```

and only the F-statistic with its associated degrees of freedom as follows:

```
summary(reg1)$fstatistic
```

```
      value      numdf      dendf  
43.09458    2.00000  29.00000
```

Exercise #3

1. Use `summary` to display the results of the regression you ran in Exercise #1 above.
2. Figure out how to extract and display *only* the regression standard error from the results of `summary` in part 1 of this exercise.
3. Calculate the regression standard error for the regression from part 1 of this exercise “by hand” and make sure that your answer matches part 2. Hint: use `resid`

Solution to Exercise #3

1. Store the result of `lm` and use `summary`:

```
myreg <- lm(mpg ~ disp + hp + cyl + wt, mtcars)  
summary(myreg)
```

Call:

```
lm(formula = mpg ~ disp + hp + cyl + wt, data = mtcars)
```

Residuals:

```
      Min       1Q   Median       3Q      Max  
-4.0562 -1.4636 -0.4281  1.2854  5.8269
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) 40.82854    2.75747  14.807 1.76e-14 ***  
disp         0.01160    0.01173   0.989 0.331386  
hp          -0.02054    0.01215  -1.691 0.102379  
cyl         -1.29332    0.65588  -1.972 0.058947 .  
wt          -3.85390    1.01547  -3.795 0.000759 ***  
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 2.513 on 27 degrees of freedom

Multiple R-squared: 0.8486, Adjusted R-squared: 0.8262

F-statistic: 37.84 on 4 and 27 DF, p-value: 1.061e-10

2. The appropriate list item is called `sigma`

```
summary(myreg)$sigma
```

```
[1] 2.51252
```

3. Let $\hat{\epsilon}_1, \dots, \hat{\epsilon}_n$ denote the residuals. Then the standard error of the regression is given by

$$\sqrt{\frac{\sum_{i=1}^n \hat{\epsilon}_i^2}{n - p - 1}}$$

where p is the number of X -variables in the regression. We can implement this in R using `resid` and compare it to the results calculated automatically by `summary` as follows:

```
ehat <- resid(myreg)
n <- length(ehat)
p <- length(coef(myreg)) - 1
sqrt(sum(ehat^2) / (n - p - 1))
```

```
[1] 2.51252
```

```
summary(myreg)$sigma
```

```
[1] 2.51252
```

Regression Without an Intercept

More than 99% of the time, it makes sense for us to include an intercept β_0 in a linear regression. To see why, consider the meaning of β_0 : this is the value of Y that we would predict if $X_1 = X_2 = \dots = X_p = 0$. Unless we have some very strong *a priori* knowledge, there is no reason to suppose that the mean of Y should be zero when all of the predictors are zero. In some very special cases, however, we *do* have such special knowledge. To *force* the intercept in a regression to be zero we use the syntax `-1`, for example

```
summary(lm(mpg ~ disp - 1, mtcars))
```

Call:

```
lm(formula = mpg ~ disp - 1, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-17.471	-2.900	7.034	14.799	29.702

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
disp	0.059049	0.009765	6.047	1.07e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 14.42 on 31 degrees of freedom

Multiple R-squared: 0.5412, Adjusted R-squared: 0.5264

F-statistic: 36.57 on 1 and 31 DF, p-value: 1.073e-06

Exercise #4

What do you get if you run the regression `lm(mpg ~ 1, mtcars)`?

Solution to Exercise #4

This calculates the sample mean of mpg

```
lm(mpg ~ 1, mtcars)
```

Call:

```
lm(formula = mpg ~ 1, data = mtcars)
```

Coefficients:

```
(Intercept)
      20.09
```

```
with(mtcars, mean(mpg))
```

```
[1] 20.09062
```

F-tests

Suppose we want to test the *joint* null hypothesis $H_0: \beta_1 = \beta_2 = \dots = \beta_p = 0$ versus the alternative that at least one of these coefficients is non-zero. This is equivalent to testing the null hypothesis that none of the predictors X_1, \dots, X_p is helpful in predicting Y . This test is automatically carried out by `summary`. Consider a regression that uses `disp`, `hp`, `wt` and `cyl` to predict `mpg`

```
reg2 <- lm(mpg ~ disp + hp + wt + cyl, mtcars)
summary(reg2)
```

Call:

```
lm(formula = mpg ~ disp + hp + wt + cyl, data = mtcars)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-4.0562 -1.4636 -0.4281  1.2854  5.8269
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  40.82854     2.75747   14.807 1.76e-14 ***
disp          0.01160     0.01173    0.989 0.331386
hp           -0.02054     0.01215   -1.691 0.102379
wt           -3.85390     1.01547   -3.795 0.000759 ***
cyl          -1.29332     0.65588   -1.972 0.058947 .
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

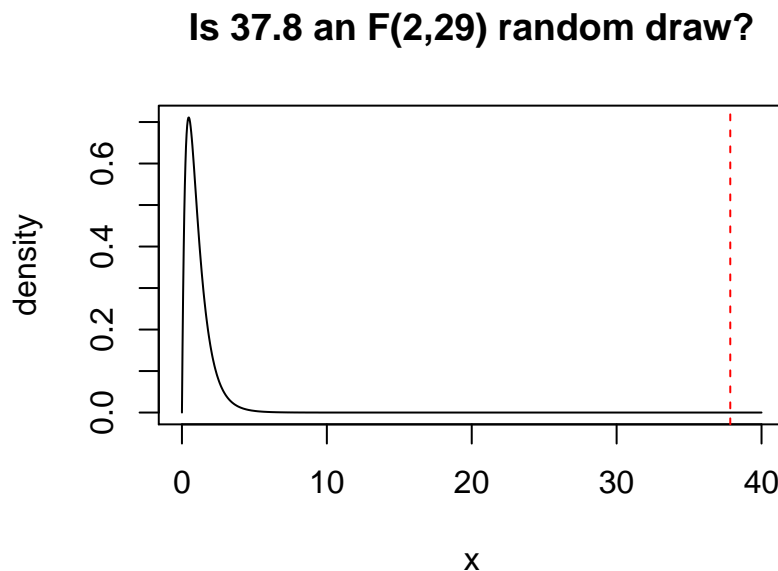
```
Residual standard error: 2.513 on 27 degrees of freedom
Multiple R-squared:  0.8486,    Adjusted R-squared:  0.8262
F-statistic: 37.84 on 4 and 27 DF,  p-value: 1.061e-10
```

At the very bottom of the `summary` output is a line that begins **F-statistic**. This line contains the results of the F-test of the joint null described above. In this case the p-value is miniscule: there is very strong evidence that at least one of the predictors is helpful in predicting `mpg`. To get a better understanding of what's involved here, we can calculate the p-value “by hand” as follows:

```
summary(reg2)$fstatistic
```

```
      value      numdf      dendf
37.84413    4.00000    27.00000
```

```
curve(expr = df(x, 4, 27), 0, 40, n = 1001,
      ylab = 'density',
      main = 'Is 37.8 an F(2,29) random draw?')
abline(v = 37.84413, lty = 2, col = 'red')
```



```
1 - pf(37.84413, 4, 27)
```

```
[1] 1.061086e-10
```

Under the null hypothesis, the F-statistic is a draw from an F random variable with numerator degrees of freedom 2 and denominator degrees of freedom 29. (If you are unfamiliar with the F-distribution see my Tutorial Friends of the Normal Distribution.) So is it plausible that the value 37.8 came from an $F(2, 29)$ distribution? From my plot of the corresponding density function, the answer is clearly *no*. We have very strong evidence against the null hypothesis.

Sometimes we only want to test the null that a *subset* of our predictors is unhelpful for predicting Y . For example, in `reg2` we might ask whether `wt` and `cyl` provide *extra* information for predicting `mpg`

after we have already included `disp` and `hp` in our model. To carry out this test, we use the function `linearHypothesis` from the package `car`. Make sure to install this package before proceeding. Note the syntax: `linearHypothesis([lm object], c('[first restriction]', ..., '[last restriction]')`

```
library(car)
```

Loading required package: carData

```
linearHypothesis(reg2, c('wt = 0', 'cyl = 0'))
```

Linear hypothesis test

Hypothesis:

`wt = 0`

`cyl = 0`

Model 1: restricted model

Model 2: `mpg ~ disp + hp + wt + cyl`

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	29	283.49				
2	27	170.44	2	113.05	8.954	0.00104 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The two key numbers to look for in the output are `F`, the value of the F-statistic, and `Pr(>F)`, the p-value. The other values are the inputs used to calculate `F`. (See Equation 3.24 in ISL.) In this instance we strongly reject the null hypothesis that `wt` and `cyl` are irrelevant for predicting `mpg` after controlling for `disp` and `hp`.

Exercise #5

Generate two vectors of independent standard normal draws `x` and `z`. Each vector should contain as many elements as there are rows in `mtcars`. Use the command `set.seed(1234)` before making your random draws so that they are replicable. (By first setting the seed to a fixed number, you ensure that the same random draws will be made any time that you re-run this code chunk.) Carry out a new regression, `reg3`, that *augments* `reg2` by adding the predictors `x` and `z`. Then carry out an F-test the null hypothesis that `x` and `z` are irrelevant for predicting `mpg` after controlling for `disp`, `hp`, `wt`, and `cyl`. Interpret your findings. Do the results of the test make sense?

Solution to Exercise #5

```
set.seed(1234)
n <- nrow(mtcars)
x <- rnorm(n)
z <- rnorm(n)
reg3 <- lm(mpg ~ disp + hp + wt + cyl + x + z, mtcars)
linearHypothesis(reg3, c('x = 0', 'z = 0'))
```

Linear hypothesis test

Hypothesis:

$x = 0$

$z = 0$

Model 1: restricted model


Model 2: $\text{mpg} \sim \text{disp} + \text{hp} + \text{wt} + \text{cyl} + x + z$

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	27	170.44				
2	25	169.95	2	0.4917	0.0362	0.9645

Plotting the Regression Line

To get an idea of whether our regression model looks reasonable, it's always a good idea to make some plots. When we have a single predictor X , it is common to plot the raw X and Y observations along with the regression line. It's easy to do this using `ggplot`. Suppose we wanted to predict `mpg` using `disp`. Here's the `ggplot` way to plot the data and regression line:

```
ggplot(mtcars, aes(x = disp, y = mpg)) +  
  geom_point() +  
  geom_smooth(method = 'lm')
```




Notice that I specified `aes` inside of `ggplot`. This ensures that both `geom_point` and `geom_smooth` “know” which variable is x and which variable is y . Notice moreover, that the `ggplot` way of doing this includes *error bounds* for the regression line. This is a handy way of visualizing the uncertainty in the line we've fit.

Exercise #6

Make a `ggplot` with `hp` on the x-axis and `mpg` on the y-axis that includes the regression line for predicting `mpg` from `hp`.

Solution to Exercise #6

```
ggplot(mtcars, aes(x = hp, y = mpg)) +  
  geom_point() +  
  geom_smooth(method = 'lm')
```



Polynomial Regression

In your next reading assignment, you'll learn about *polynomial regression*. The “linear” in linear regression does not actually refer to the relationship between Y and the predictors X ; it refers to the relationship between Y and the *coefficients* $\beta_0, \beta_1, \dots, \beta_p$. In the expression $Y = \beta_0 + \beta_1 X + \epsilon$, Y is a linear function of β_0 and β_1 and it is *also* a linear function of X . In the expression $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \epsilon$, Y is *still* a linear function of the coefficients, but a *quadratic* function of X . This is a simple example of polynomial regression, which allows us to model more complicated relationships between X and Y . Notice, for example, that the relationship between `mpg` and `disp` looks like it might be curved. To accommodate such a relationship, let's try a polynomial regression that includes `disp` and `disp^2`. To do this we use the syntax `I([some transformation of a predictor])`

```
reg3 <- lm(mpg ~ disp + I(disp^2), mtcars)  
summary(reg3)
```

```
Call:
lm(formula = mpg ~ disp + I(disp^2), data = mtcars)


Residuals:
    Min       1Q   Median       3Q      Max
-3.9112 -1.5269 -0.3124  1.3489  5.3946

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.583e+01  2.209e+00  16.221 4.39e-16 ***
disp        -1.053e-01  2.028e-02  -5.192 1.49e-05 ***
I(disp^2)     1.255e-04  3.891e-05   3.226  0.0031 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.837 on 29 degrees of freedom
Multiple R-squared:  0.7927,    Adjusted R-squared:  0.7784
F-statistic: 55.46 on 2 and 29 DF,  p-value: 1.229e-10
```

Notice that the coefficient on the quadratic term is *highly* statistically significant, which is strong evidence of curvature in the relationship between `mpg` and `disp`. We can plot the polynomial regression as follows:

```
ggplot(mtcars, aes(x = disp, y = mpg)) +
  geom_point() +
  geom_smooth(method = 'lm', formula = y ~ x + I(x^2))
```




Notice that this requires us to specify the `formula` argument so that `ggplot` knows that we want to plot a quadratic relationship.

Exercise #7

In my code above, I considered a quadratic relationship between `mpg` and `disp`. Add a cubic term to the regression, plot the points and regression function, and display the results using `summary`. Comment on the results.

Solution to Exercise #7

```
reg4 <- lm(mpg ~ disp + I(disp^2) + I(disp^3), mtcars)
ggplot(mtcars, aes(x = disp, y = mpg)) +
  geom_point() +
  geom_smooth(method = 'lm', formula = y ~ x + I(x^2) + I(x^3))
```



```
summary(reg4)
```

Call:

```
lm(formula = mpg ~ disp + I(disp^2) + I(disp^3), data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.0896	-1.5653	-0.3619	1.4368	4.7617

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	5.070e+01	3.809e+00	13.310	1.25e-13	***
disp	-3.372e-01	5.526e-02	-6.102	1.39e-06	***
I(disp^2)	1.109e-03	2.265e-04	4.897	3.68e-05	***

```

I(displ^3)    -1.217e-06  2.776e-07  -4.382  0.00015 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.224 on 28 degrees of freedom
Multiple R-squared:  0.8771,    Adjusted R-squared:  0.8639 
F-statistic: 66.58 on 3 and 28 DF,  p-value: 7.347e-13

```

Interaction Effects

An idea closely related to polynomial regression that will also be discussed in your next reading assignment is that of an *interaction*. In the model $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1 X_2 + \epsilon$, Y is a linear function of $\beta_0, \beta_1, \beta_2$, and β_3 but a *nonlinear* function of X_1 and X_2 . The term $X_1 \times X_2$ is called an *interaction*. To run a regression with an interaction, we use the syntax `[One Predictor]:[Another Predictor]` for example

```
lm(mpg ~ displ + hp + displ:hp, mtcars)
```

Call:

```
lm(formula = mpg ~ displ + hp + displ:hp, data = mtcars)
```

Coefficients:

(Intercept)	displ	hp	displ:hp
39.67426	-0.07337	-0.09789	0.00029

Exercise #8

Fit a regression using `displ`, `displ^2`, `wt`, `wt^2` and the interaction between `wt` and `displ` to predict `mpg` and display the coefficient estimates.

Solution to Exercise #8

```
reg5 <- lm(mpg ~ displ + wt + I(displ^2) + I(wt^2) + displ:wt, mtcars)
coef(reg5)
```

(Intercept)	displ	wt	I(displ^2)	I(wt^2)
4.692786e+01	-3.172401e-02	-1.062827e+01	2.019044e-04	2.079131e+00
displ:wt				
-2.660633e-02				

Predicting New Observations

To predict new observations based on a fitted linear regression model in R, we use the `predict` function. For example, consider three hypothetical cars with the following values of displacement and horsepower


```
mydat <- data.frame(displacement = c(100, 200, 300),
                    horsepower = c(150, 100, 200))
mydat
```

```
  displacement horsepower
1      100      150
2      200      100
3      300      200
```

Based on the results of `reg1`, we would predict that these cars have the following values of `mpg`

```
predict(reg1, mydat)
```

```
      1      2      3
23.97526 22.18264 16.66401
```

Note the syntax: the first argument of `predict` is a set of regression results while the second is a dataframe (or tibble) with column names that *match* the variables in the regression we carried out.

Exercise #9

1. Check the predictions of `predict` in the preceding chunk “by hand” using `mydat`, `coef`, and `reg1`.
2. Consider three cars with `displacement` equal to 125, 175, and 225, respectively. Predict `mpg` for each of these based on the regression from Exercise #7.

Solution to Exercise #9

1. Many possibilities. Here’s one:

```
b <- coef(reg1)
b0 <- b[1]
b1 <- b[2]
b2 <- b[3]
b0 + b1 * mydat$displacement + b2 * mydat$horsepower
```

```
[1] 23.97526 22.18264 16.66401
```

2. Use the following:

```
mydat <- data.frame(displacement = c(125, 175, 225))
predict(reg1, mydat)
```

```
      1      2      3
23.50258 19.13630 17.12228
```