Lab #9 - Logistic Regression Part I

Econ 224

September 25th, 2018

Notation

In this lab we'll be looking at logistic regression, a commonly-used binary classification method. To make things simpler, we'll use some slightly different notation and terminology than ISL. First we'll define the *column vectors* X and β as follows:

$$X = \begin{bmatrix} 1 \\ X_1 \\ X_2 \\ \vdots \\ X_p \end{bmatrix}, \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix}$$

Notice that the first element of X is *not* X_1: it is simply the number 1. There's an important reason for this that you'll see in a moment. From the reading, we know that logistic regression is a *linear model* for the *log odds*, namely

$$\log\left[\frac{P(X)}{1-P(X)}\right] = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$$

where $P(X)$ is shorthand for $\mathbb{P}(Y = 1|X)$. Note that when I write log I **always** mean the natural logarithm. Also note that when I write $\exp(z)$ I mean e^z. This comes in handy if z is a complicated expression.

Using the vector notation introduced above, we can express this more compactly as

$$\log\left[\frac{P(X)}{1-P(X)}\right] = X'\beta$$

since

$$X'\beta = \begin{bmatrix} 1 & X_1 & X_2 & \cdots & X_p \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix} = \beta_0 + \beta_1 X_2 + \beta_2 X_2 + \cdots + \beta_p X_p$$

I will call $X'\beta$ the *linear predictor* since it is the linear function of X that we use to predict Y. By exponentiating both sides of the log-odds expression from above and re-arranging, obtain the following:

$$\begin{aligned} \frac{P(X)}{1-P(X)} &= \exp(X'\beta) \\ P(X) &= [1-P(X)]\exp(X'\beta) \\ P(X) + P(X)\exp(X'\beta) &= \exp(X'\beta) \\ P(X)[1+\exp(X'\beta)] &= \exp(X'\beta) \\ P(X) &= \frac{\exp(X'\beta)}{1+\exp(X'\beta)} \\ P(X) &= \Lambda(X'\beta) \end{aligned}$$

where the function Λ is defined as follows

$$\Lambda(z) = \frac{e^z}{1+e^z}$$

Exercise #1

(a) Verify that $\Lambda(z) = \dfrac{1}{1+e^{-z}}$.
(b) Using (b), write an alternative expression for $P(X)$.

Solution to Exercise #1

(a) Dividing the numerator and denominator by e^z, which cannot result in division by zero since e^z is always positive, we have

$$\Lambda(z) = \frac{e^z}{1+e^z} = \frac{1}{1/e^z + 1} = \frac{1}{1+e^{-z}}$$

(b)

Interpreting β in a Logistic Regression

From the expression above, we see that β_j is the partial derivative of the log-odds with respect to X_j. But it's difficult to think in terms of log-odds. By doing some calculus (see the exercises below), we can work out the partial derivative of $p(X)$ with respect to X_j, but this will *not* turn out to equal β_j. Because $P(X)$ is not a linear function of X, the derivative varies with X, which makes things fairly complicated. There are two main approaches for dealing with this problem. One is to evaluate the derivative at a "typical" value of X such as the sample mean. Another is to use the "divide by 4 rule." This rule says that if we increase X_j by one unit, $P(X)$ will change by *no more than* $\beta_j/4$. In the following exercise, you'll derive this rule.

Exercise #2

(a) Analyze the function $\Lambda(z)$: calculate its derivative, and its limits as $z \to -\infty$ and $+\infty$. What values can this function take? Is it increasing? Decreasing? Explain.
(b) Use the chain rule and your answer to (a) to find the partial derivative of $\Lambda(X'\beta)$ with respect to X_j.
(c) What is the maximum value of the *derivative* of $\Lambda(z)$? At what value of z does it occur?
(d) Use your answers to parts (a), (b) and (c) to justify the "divide by 4 rule."
(e) The "divide by 4 rule" provides an upper bound on the effect of X_j on $P(X)$. When is this upper bound close to the derivative you calculated in part (c)?

Solution to Exercise #2

(a) The function Λ takes values between 0 and 1. When $z = 0$, $\Lambda(z) = e^0/(1+e^0) = 1/2$. As $z \to \infty$, $\Lambda(z) \to 1$ and as $z \to -\infty$, $\Lambda(z) \to 0$. We calculate its derivative using the quotient rule as follows

$$\frac{d\Lambda(z)}{dz} = \frac{e^z(1+e^z) - e^z e^z}{(1+e^z)^2} = \frac{e^z}{(1+e^z)^2}$$

Since e^z is always greater than zero, the derivative is always positive so $\Lambda(z)$ is strictly increasing.
(b) The key is to treat the linear predictor $X'\beta$ as a function of X_j, namely

$$f(X_j) = X'\beta = \beta_0 + \beta_1 X_1 + \cdots + \beta_j X_j + \beta_{j+1} X_{j+1} + \cdots + \beta_p X_p$$

Now, by the chain rule we have

$$\frac{\partial \Lambda(X'\beta)}{\partial X_j} = \frac{\partial \Lambda\left(f(X_j)\right)}{\partial X_j} = \frac{\exp(X'\beta)}{[1+\exp(X'\beta)]^2}\frac{\partial f(X_j)}{\partial X_j} = \frac{\beta_j \exp(X'\beta)}{[1+\exp(X'\beta)]^2}$$

(c) To find the value of z that maximizes the first derivative, we take the *second* derivative of Λ as follows

$$\begin{aligned}\frac{d^2\Lambda(z)}{dz} &= \frac{e^z(1+e^z)^2 - 2e^z(1+e^z)e^z}{(1+e^z)^4} = \frac{e^z(1+2e^z+e^{2z}) - 2e^{2z}(1+e^z)}{(1+e^z)^4} \\ &= \frac{e^z + 2e^{2z} + e^{3z} - 2e^{2z} - 2e^{3z}}{(1+e^z)^4} = \frac{e^z - e^{3z}}{(1+e^z)^4} \\ &= \frac{e^z(1-e^{2z})}{(1+e^z)^4} = \frac{e^z(1+e^z)(1-e^z)}{(1+e^z)^4} = \frac{e^z(1-e^z)}{(1+e^z)^3}\end{aligned}$$

Thus, the first order condition is $e^z(1-e^z) = 0$. Since e^z cannot equal zero for any z, the only way for this equation to be satisfied is if $e^z = 1$ which occurs precisely when $z = 0$. Substituting into our expression from (a), we find that the derivative of $\Lambda(z)$ at $z = 0$ is $e^0/(1+e^0)^2 = 1/(1+1)^2 = 1/4$.

(d) From part (a), we know that the derivative of $\Lambda(z)$ equals $e^z/(1+e^z)^2$ which is always positive. From part (c) we know that this derivative is *at most* $1/4$. Therefore, the partial derivative of $\Lambda(X'\beta)$ with respect to X_j is *at most* $\beta_j \times 1/4 = \beta_j/4$.

(e) When $X'\beta \approx 0$ it follows that $\exp(X'\beta)/[1+\exp(X'\beta)]^2 \approx 1/4$ so the "divide by four" rule gives a good approximation to the actual derivative.

Contaminated Wells in Bangladesh

In the remainder of this lab, as well as Thursday's lab, we'll work with a dataset containing household-level information from Bangladesh: `wells.csv`. You can download the dataset from the course website at http://ditraglia.com/econ224/wells.csv.

Here is some background on the dataset from Gelman and Hill (2007):

> Many of the wells used for drinking water in Bangladesh and other South Asian countries are contaminated with natural arsenic ... a research team from the United States and Bangladesh measured all the wells [in a small region] and labeled them with their arsenic level as well as a characterization of "safe" (below 0.5 in units of hundreds of micrograms per liter, the Bangladesh standard for arsenic in drinking water) or "unsafe" (above 0.5). People with unsafe wells were encourage to switch to nearby private or community wells or to new wells of their own construction. A few years later, the researchers returned to find out who had switched wells.

Our goal is to predict which households will switch wells using the following information:

Name	Description
`dist`	Distance to closest known safe well (meters)
`arsenic`	Arsenic level of respondent's well (100s of micrograms/liter)
`switch`	Dummy variable: equals 1 if switched to a new well
`assoc`	Dummy variable: equals 1 if any member of the household is active in community organizations
`educ`	Education level of head of household (years)

Running a Logistic Regression in R